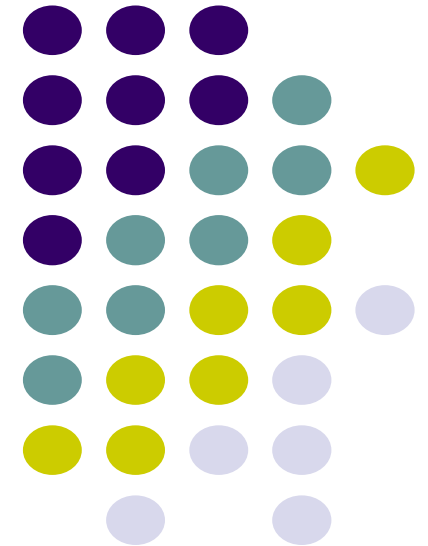


The Basics of UNIX/Linux

11-1. Arrays and Pointer. Part 1

Instructor: Joonho Kwon
jhkwon@pusan.ac.kr
Data Science Lab @ PNU



Lecture Outline




- **C's Memory Model** (refresher)
- Pointers (refresher)
- Arrays

OS and Processes




- The OS lets you run multiple applications at once
 - An application runs within an OS “process”
 - The OS timeslices each CPU between runnable processes
 - This happens *very quickly*: ~100 times per second



Processes and Virtual Memory



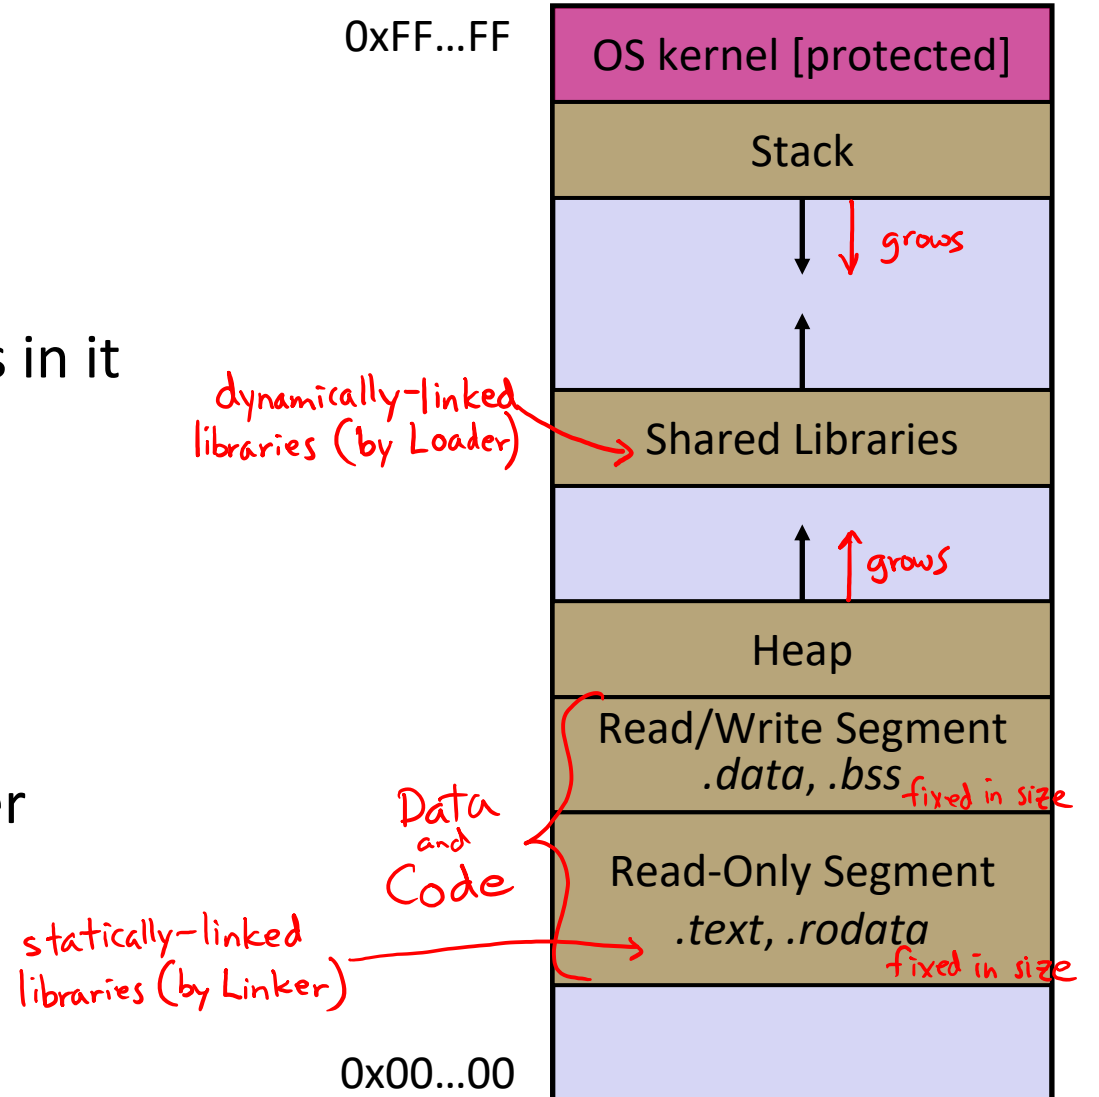
- The OS gives each process the illusion of its own private memory
 - Called the process' **address space**
 - Contains the process' virtual memory, visible only to it (via translation)
 - 2^{64} bytes on a 64-bit machine



Loading




- When the OS loads a program it:
 - 1) Creates an address space
 - 2) Inspects the executable file to see what's in it
 - 3) (Lazily) copies regions of the file into the right place in the address space
 - 4) Does any final linking, relocation, or other needed preparation



Memory Management



- *Local* variables on the Stack
 - Allocated and freed via calling conventions (`push`, `pop`, `mov`)
- *Global* and *static* variables in Data
 - Allocated/freed when the process starts/exits
- *Dynamically-allocated* data on the Heap
 - `malloc()` to request; `free()` to free, otherwise **memory leak**



Stack in Action (1/4)

Note: arrow points to *next* instruction to be executed (like in gdb).



stack.c

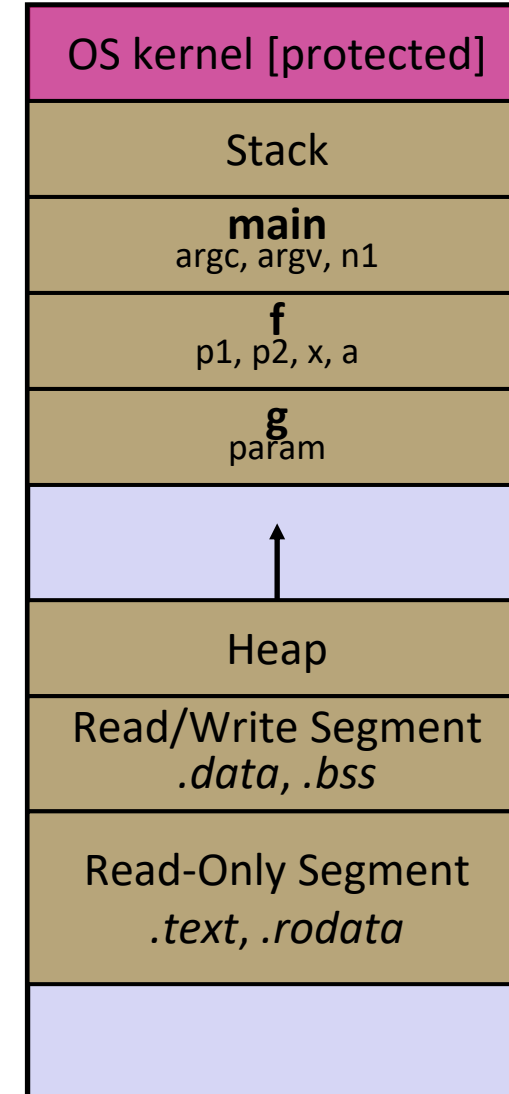
```
#include <stdint.h>

int f(int, int);
int g(int);

→ int main(int argc, char** argv) {
→   int n1 = f(3, -5);
    n1 = g(n1);
}

→ int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

→ int g(int param) {
→   return param * 2;
}
```



Stack in Action (2/4)

Note: arrow points to *next* instruction to be executed (like in gdb).



stack.c

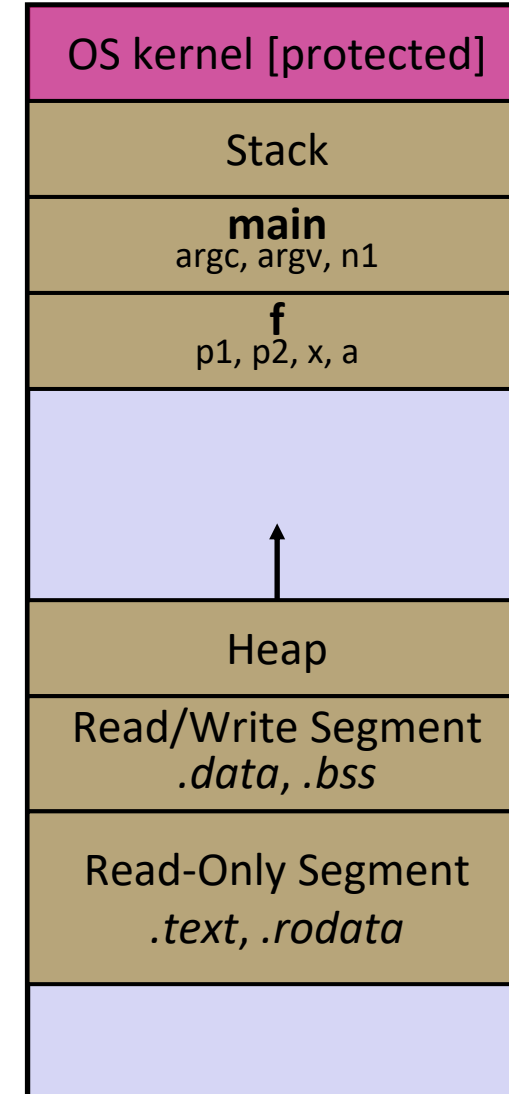
```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Stack in Action (3/4)

Note: arrow points to *next* instruction to be executed (like in gdb).



```
#include <stdint.h>

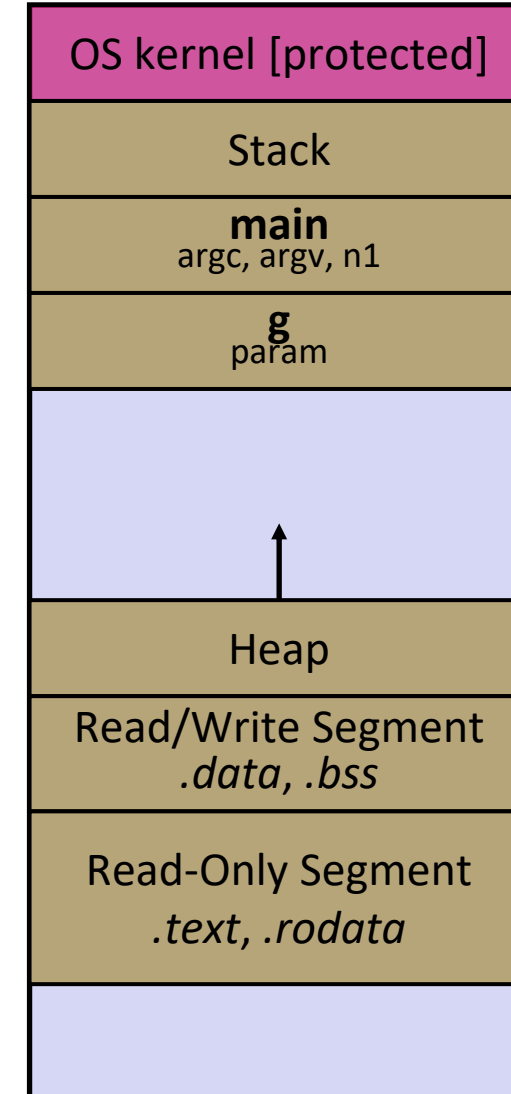
int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```

Red arrows point to the first and last lines of the `main` function.



Stack in Action (4/4)

Note: arrow points to *next* instruction to be executed (like in gdb).



stack.c


```
#include <stdint.h>

int f(int, int);
int g(int);

int main(int argc, char** argv) {
    int n1 = f(3, -5);
    n1 = g(n1);
}

int f(int p1, int p2) {
    int x;
    int a[3];
    ...
    x = g(a[2]);
    return x;
}

int g(int param) {
    return param * 2;
}
```



Lecture Outline



- C's Memory Model (refresher)
- **Pointers** (refresher)
- Arrays

Pointers



- Variables that store addresses
 - It points to somewhere in the process' virtual address space
 - `&foo` produces the virtual address of `foo`

- Generic definition: `type* name;` `type *name;`

- Recommended to not define multiple pointers on same line:

```
int *p1, p2;
```

not the same as

```
int *p1, *p2;
```

- Instead, use:

```
int *p1;  
int *p2;
```

- *Dereference* a pointer using the unary `*` operator
 - Access the memory referred to by a pointer

Pointer Example



pointy.c

```
#include <stdio.h>
#include <stdint.h>

int main(int argc, char** argv) {
    int x = 351;
    int* p;      // p is a pointer to a int

    p = &x;      // p now contains the addr of x
    printf("&x is %p\n", &x);
    printf(" p is %p\n", p);
    printf(" x is %d\n", x);

    *p = 333;    // change value of x
    printf(" x is %d\n", x);

    return 0;
}
```

Something Curious



- What happens if we run `pointy.c` several times?

```
$ gcc -Wall -std=c11 -o pointy pointy.c
```

Run 1: **bash\$./pointy**
`&x is 0x7fff`**f9e28524**
`p is 0x7ffff9e28524`
`x is 351`
`x is 333`

Run 2: **bash\$./pointy**
`&x is 0x7fff`**e847be34**
`p is 0x7fffe847be34`
`x is 351`
`x is 333`


Run 3: **bash\$./pointy**
`&x is 0x7fff`**e7b14644**
`p is 0x7fffe7b14644`
`x is 351`
`x is 333`

Run 4: **bash\$./pointy**
`&x is 0x7fff`**ff0dfe54**
`p is 0x7fffff0dfe54`
`x is 351`
`x is 333`



Address Space Layout Randomization

- Linux uses *address space layout randomization* (ASLR) for added security
 - Randomizes:
 - Base of stack
 - Shared library (`mmap`) location
 - Makes Stack-based buffer overflow attacks tougher
 - Makes debugging tougher
 - Can be disabled (`gdb` does this by default); Google if curious



Lecture Outline



- C's Memory Model (refresher)
- Pointers (refresher)
- **Arrays**

Arrays



- Definition:

```
type name [size]
```

- Allocates $\text{size} * \text{sizeof}(\text{type})$ bytes of *contiguous* memory
- Normal usage is a **compile-time constant** for `size` (e.g. `int scores[175];`)
- Initially, array values are “garbage”

Size of an array



- Size of an array
 - Not stored anywhere – array does not know its own size!
 - `sizeof(array)` only works in variable scope of array definition
 - Recent versions of C allow for variable-length arrays
 - Uncommon and can be considered bad practice [*we won't use*]
 - <https://gcc.gnu.org/onlinedocs/gcc/Variable-Length.html>

```
int n = 175;  
int scores[n]; // OK in C99
```

Challenge Question



should malloc
instead of using
vla's!

- The code snippets both use a variable-length array. What will happen when we compile with C99?

allocated in
Static Data
(can't change
size)

```
int m = 175;
int scores[m];

void foo(int n) {
    ...
}
```



```
int m = 175;

void foo(int n) {
    int scores[n];
    ...
}
```

✓ allocated on
the stack
(can grow)

however, you
don't want to
put large arrays
on the stack

- | | | |
|----|----------------|----------------|
| A. | Compiler Error | Compiler Error |
| B. | Compiler Error | No Error |
| C. | No Error | Compiler Error |
| D. | No Error | No Error |
| E. | We're lost... | |

Using Arrays (1/2)



- Initialization

```
type name[size] = {val0, ..., valN};
```

- { } initialization can *only* be used at time of definition
- If no **size** supplied, infers from length of array initializer

Using Arrays (2/2)



- Array name used as identifier for “collection of data”
 - `name[index]` specifies an element of the array and can be used as an assignment target or as a value in an expression
 - **Array name** (by itself) produces **the address of the start of the array**
 - Cannot be assigned to / changed

```
int primes[6] = {2, 3, 5, 6, 11, 13};  
primes[3] = 7;  
primes[100] = 0; // memory smash!
```

Multi-dimensional Arrays



- Generic 2D format:

```
type name[rows][cols] ={{values}, ..., {values}};
```

- Still allocates a single, contiguous chunk of memory
- C is *row-major*

```
// a 2-row, 3-column array of doubles
double grid[2][3];

// a 3-row, 5-column array of ints
int matrix[3][5] = {
    {0, 1, 2, 3, 4},
    {0, 2, 4, 6, 8},
    {1, 3, 5, 7, 9}
};
```

Parameters: reference vs. value (1/2)



- There are two fundamental parameter-passing schemes in programming languages
- **Call-by-value / Pass by value**
 - Parameter is a local variable initialized when the function is called and **gets a copy of the calling argument**
 - manipulating the parameter only changes copy, *not* the calling argument
 - **C, Java, C++** primitives

Parameters: reference vs. value (2/2)



- Call-by-reference / Pass-by-reference
 - Parameter is **an alias** for the supplied argument
 - manipulating the parameter manipulates the calling argument
 - C++ references (you'll learn later)

Arrays as Parameters



- It's tricky to use arrays as parameters
 - What happens when you use an array name as an argument?
 - Arrays do not know their own size

```
int sumAll(int a[]); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers);
    return 0;
}

int sumAll(int a[]) {
    int i, sum = 0;
    for (i = 0; i < ...???)
}
```

Solution 1: Declare Array Size



```
int sumAll(int a[5]); // prototype

int main(int argc, char** argv) {
    int numbers[] = {9, 8, 1, 9, 5};
    int sum = sumAll(numbers);
    printf("sum is: %d\n", sum);
    return 0;
}

int sumAll(int a[5]) {
    int i, sum = 0;
    for (i = 0; i < 5; i++) {
        sum += a[i];
    }
    return sum;
}
```

- Problem: loss of generality/flexibility!

Solution 2: Pass Size as Parameter



- This is the standard idiom in C programs

```
int sumAll(int a[], int size); //  
prototype  
  
int main(int argc, char** argv) {  
    int numbers[] = {9, 8, 1, 9, 5};  
    int sum = sumAll(numbers, 5);  
    printf("sum is: %d\n", sum);  
    return 0;  
}  
  
int sumAll(int a[], int size) {  
    int i, sum = 0;  
    for (i = 0; i < size; i++) {  
        sum += a[i];  
    }  
    return sum;  
}
```

Returning an Array



- Local variables, including arrays, are allocated on the Stack
 - They “disappear” when a function returns!
 - Can’t safely return local arrays from functions
 - Can’t return an array as a return value – why not?

```
int* copyArray(int src[], int size) {  
    int i, dst[size];    // OK in C99  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
  
    return dst;    // no compiler error, but wrong!  
}
```

buggy_copyarray.c

Solution: Output Parameter



- Create the “returned” array in the caller
 - Pass it as an **output parameter** to `copyarray()`
 - A pointer parameter that allows the callee to leave values for the caller to use
 - Works because arrays are “passed” as pointers
 - “Feels” like call-by-reference, *but it's not*

```
void copyArray(int src[], int dst[], int size) {  
    int i;  
  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];  
    }  
}
```

copyarray.c

Arrays: Call-By-Value or Call-By-Reference?



- **Technical answer:**

- a $T[]$ array parameter is “promoted” to a pointer of type T^* , and the *pointer* is passed by value
 - So it acts like a call-by-reference array (if callee changes the array parameter elements it changes the caller’s array)
 - But it’s really a call-by-value pointer (the callee can change the pointer parameter to point to something else(!))

```
void copyArray(int src[], int dst[], int size) {  
    int i;  
    dst = src;  
    for (i = 0; i < size; i++) {  
        dst[i] = src[i];    // copies source array to itself!  
    }  
}
```

Output Parameters



- Output parameters are common in library functions

- `long int strtol(char* str, char** endptr, int base);`
- `int sscanf(char* str, char* format, ...);`

```
int    num, i;
char*  pEnd, str1 = "333 rocks";
char   str2[10];

// converts "333 rocks" into long -- pEnd is
// conversion end
num = (int) strtol(str1, &pEnd, 10);

// reads string into arguments based on format string
num = sscanf("3 blind mice", "%d %s", &i, str2);
```

outparam.c

Q&A

