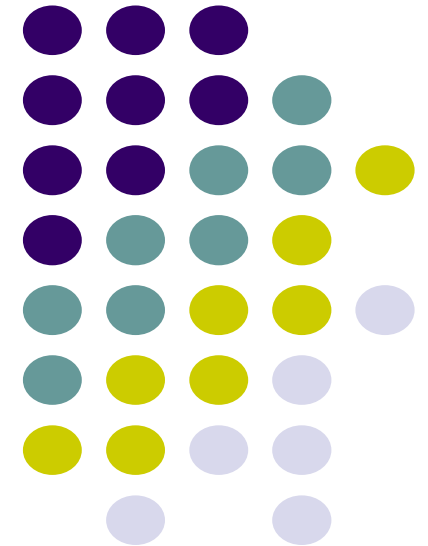


# The Basics of UNIX/Linux

## 12-1. Dynamic Memory Allocation

Instructor: Joonho Kwon  
jhkwon@pusan.ac.kr  
Data Science Lab @ PNU



# Lecture Outline



- **Heap-allocated Memory**
  - `malloc()` and `free()`
  - Memory leaks
  - Sample codes

# Memory Allocation So Far (1/2)



- So far, we have seen two kinds of memory allocation:

```
int counter = 0;    // global var

int main(int argc, char** argv) {
    counter++;
    printf("count = %d\n", counter);
    return 0;
}
```

- counter is **statically**-allocated
  - Allocated when program is loaded
  - Deallocated when program exits

# Memory Allocation So Far (2/2)



- So far, we have seen two kinds of memory allocation:

```
int foo(int a) {  
    int x = a + 1;    // local var  
    return x;  
}  
  
int main(int argc, char** argv) {  
    int y = foo(10);  // local var  
    printf("y = %d\n", y);  
    return 0;  
}
```

- $a$ ,  $x$ ,  $y$  are **automatically**-allocated
  - Allocated when function is called
  - Deallocated when function returns

# Dynamic Allocation



- Situations where static and automatic allocation aren't sufficient:
  - We need memory that persists across multiple function calls but not the whole lifetime of the program
  - We need more memory than can fit on the Stack
  - We need memory whose size is not known in advance to the caller

```
// this is pseudo-C code  
char* ReadFile(char* filename) {  
    int size = GetFileSize(filename);  
    char* buffer = AllocateMem(size);  
  
    ReadFileIntoBuffer(filename, buffer);  
    return buffer;  
}
```

# Dynamic Allocation



- What we want is *dynamically*-allocated memory
  - Your program explicitly requests a new block of memory
    - The language allocates it at runtime, perhaps with help from OS
  - Dynamically-allocated memory persists until either:
    - Your code explicitly deallocated it (manual memory management)
    - A garbage collector collects it (automatic memory management)
- C requires you to manually manage memory
  - Gives you more control, but causes headaches

# Aside: NULL



- NULL is a memory location that is **guaranteed to be invalid**
  - In C on Linux, NULL is 0x0 and an attempt to dereference NULL *causes a segmentation fault*
- Useful as an indicator of an uninitialized (or currently unused) pointer or allocation error
  - It's better to cause a segfault than to allow the corruption of memory!

segfault.c

```
int main(int argc, char** argv) {  
    int* p = NULL;  
    *p = 1; // causes a segmentation fault  
    return 0;  
}
```

# malloc()



- General usage:

```
var = (type*) malloc(size in bytes)
```

- **malloc** allocates a block of memory of the requested size
  - Returns a pointer to the first byte of that memory
    - And **returns NULL** if the memory allocation failed!
  - You should assume that the memory initially contains garbage
  - You'll typically use **sizeof** to calculate the size you need

```
// allocate a 10-float array
float* arr = (float*) malloc(10*sizeof(float));
if (arr == NULL) {
    return errcode;
}
... // do stuff with arr
```



# calloc()



- General usage:

```
var = (type*) calloc(num, bytes per element)
```

- Like **malloc**, but also zeros out the block of memory
  - Helpful for shaking out bugs
  - Slightly slower; preferred for non-performance-critical code
  - **malloc** and **calloc** are found in `stdlib.h`

```
// allocate a 10-double array
double* arr = (double*) calloc(10, sizeof(double));
if (arr == NULL) {
    return errcode;
}
...    // do stuff with arr
```

# free()



- Usage:

```
free(pointer);
```


- Deallocates the memory pointed-to by the pointer
  - Pointer *must* point to the first byte of heap-allocated memory (*i.e.* something previously returned by **malloc** or **calloc**)
  - Freed memory becomes eligible for future allocation
  - Pointer is unaffected by call to free
    - Defensive programming: can set pointer to NULL after freeing it

```
float* arr = (float*) malloc(10*sizeof(float));  
if (arr == NULL)  
    return errcode;  
...           // do stuff with arr  
free(arr);  
arr = NULL;   // OPTIONAL
```

# The Heap



- The Heap is a large pool of unused memory that is used for dynamically-allocated data
  - **malloc** allocates chunks of data in the Heap; **free** deallocates those chunks
  - **malloc** maintains bookkeeping data in the Heap to track allocated blocks



# Heap and Stack Example (1/11)



arraycopy.c

```
#include <stdlib.h>

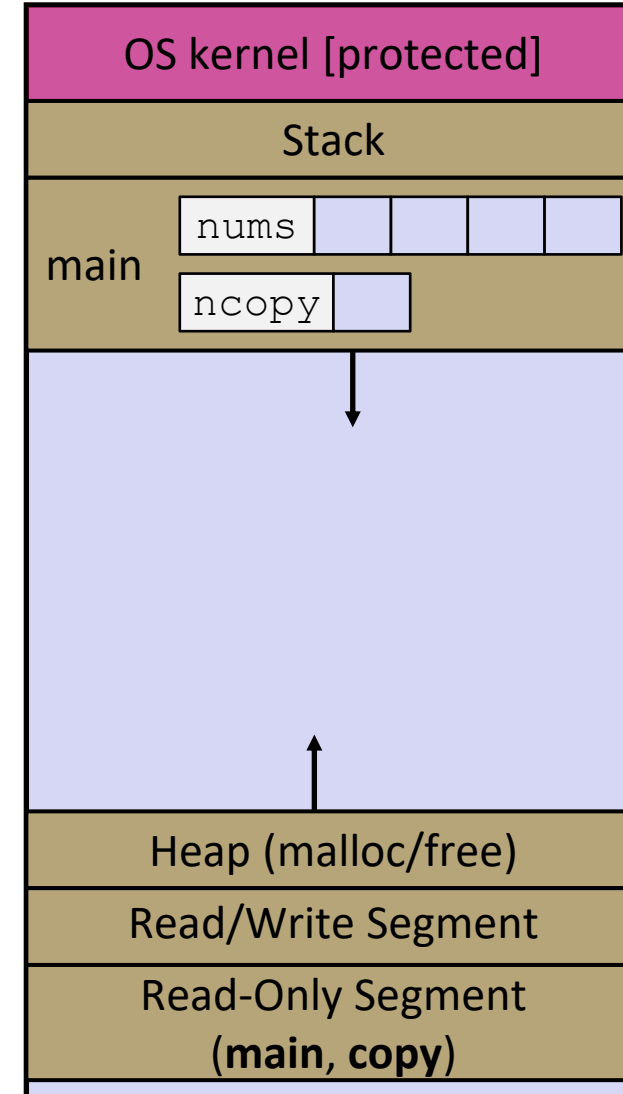
int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (2/11)



arraycopy.c

```
#include <stdlib.h>

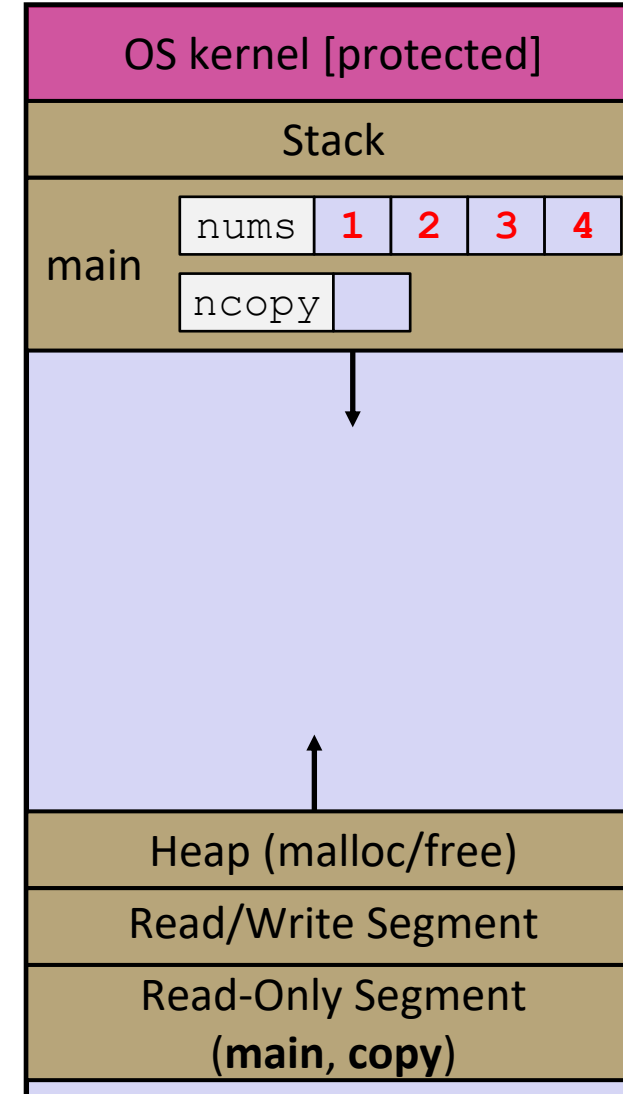
int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (3/11)



arraycopy.c

```
#include <stdlib.h>

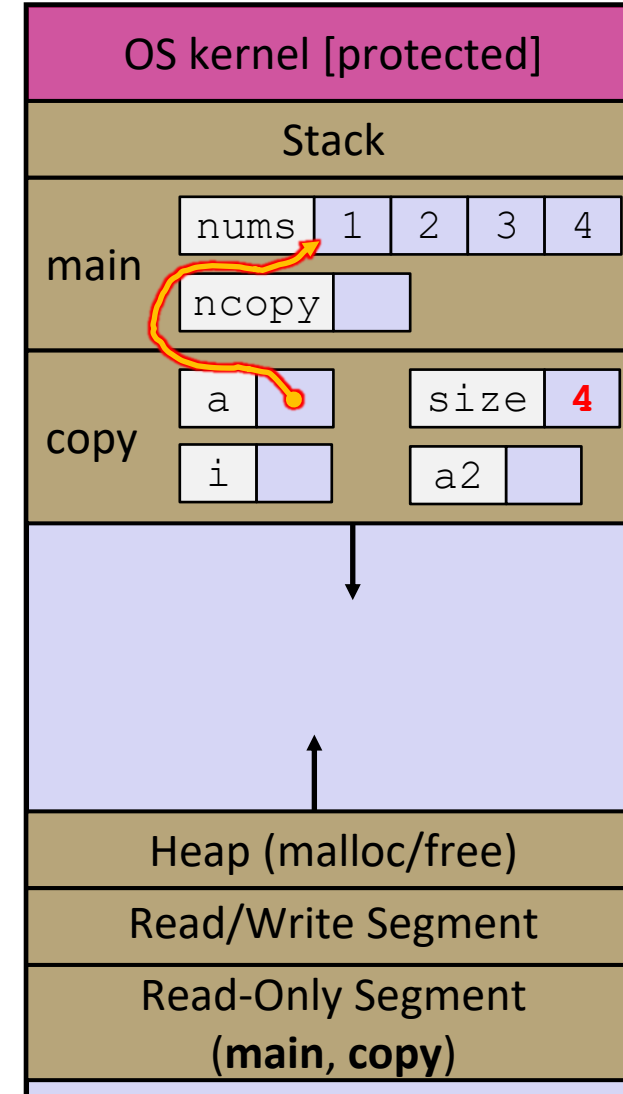
int* copy(int a[], int size) {
    int i, *a2;

    → a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (4/11)



arraycopy.c

```
#include <stdlib.h>

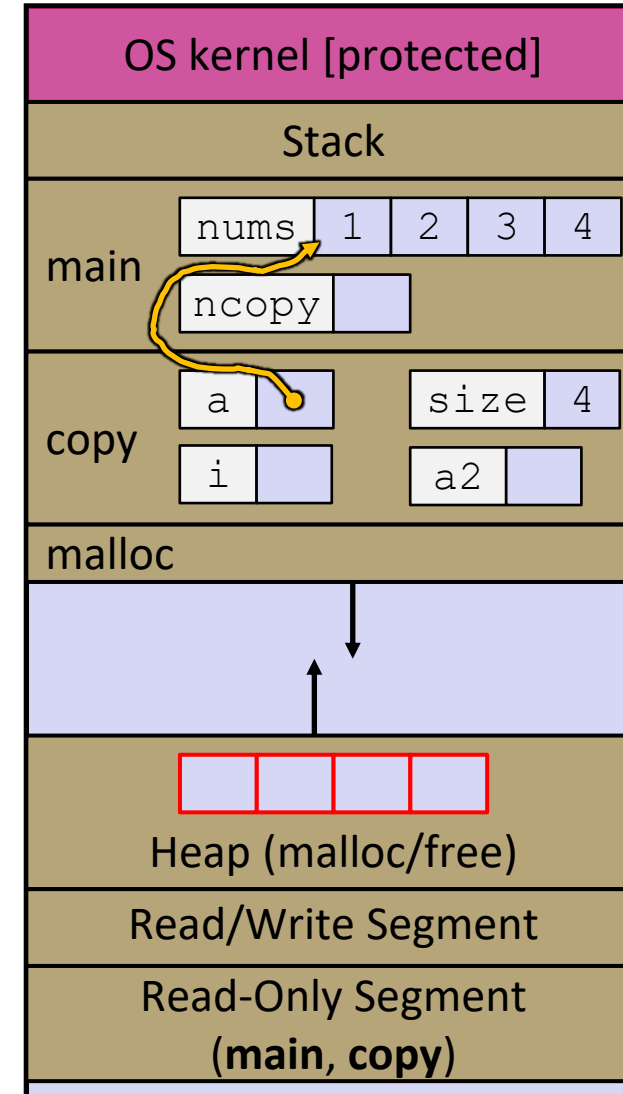
int* copy(int a[], int size) {
    int i, *a2;

    → a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (5/11)



arraycopy.c

```
#include <stdlib.h>

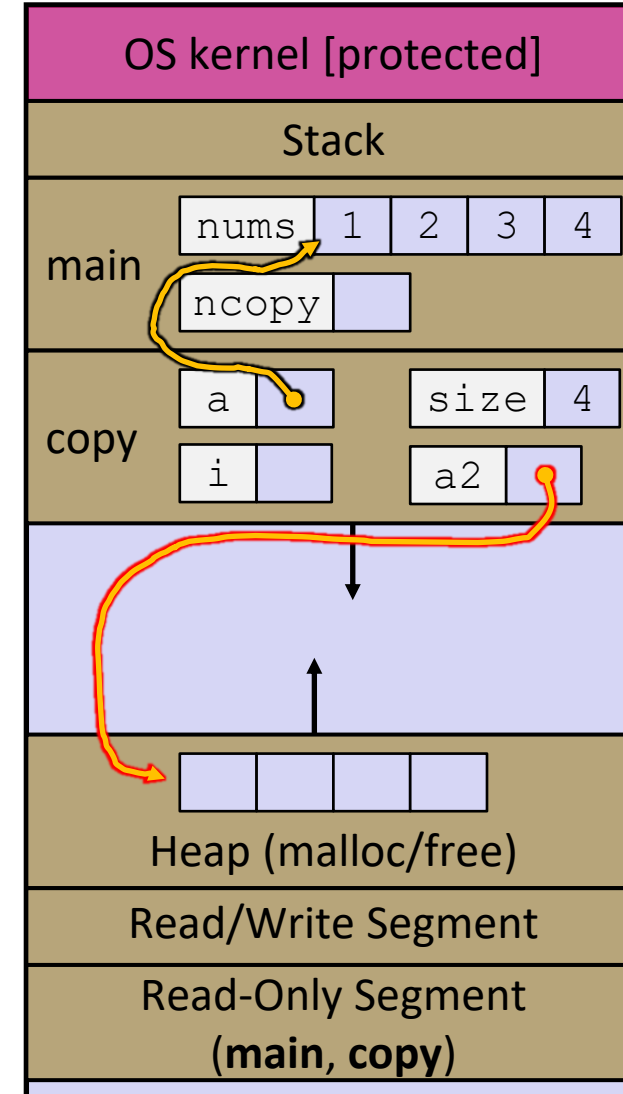
int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```





# Heap and Stack Example (6/11)



```
#include <stdlib.h>

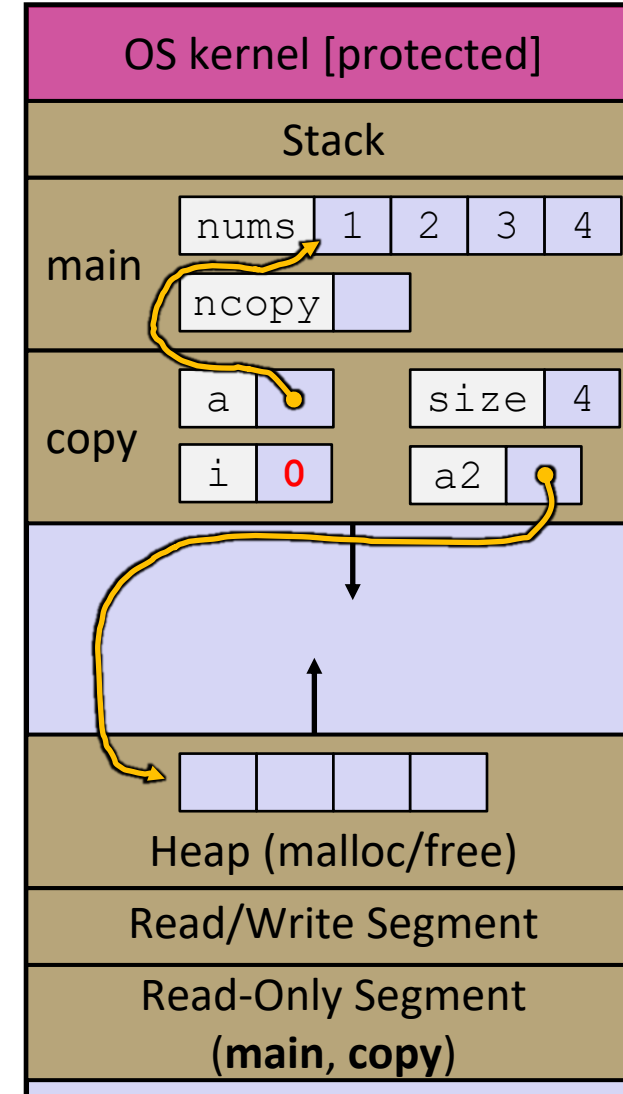
int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (7/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (8/11)



arraycopy.c

```
#include <stdlib.h>

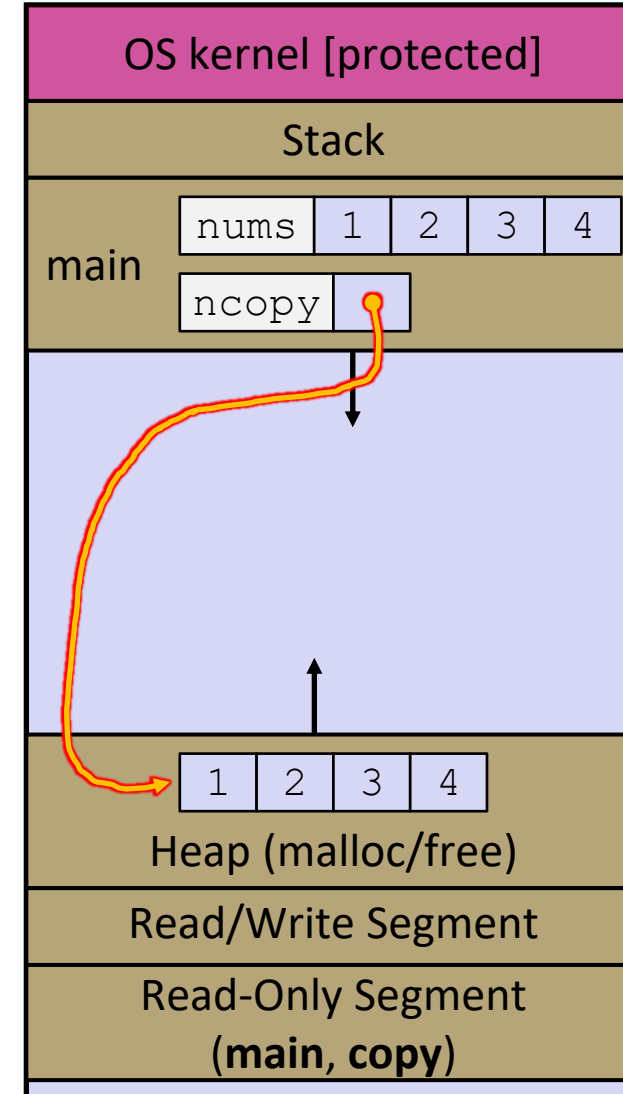
int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (9/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (10/11)



arraycopy.c

```
#include <stdlib.h>

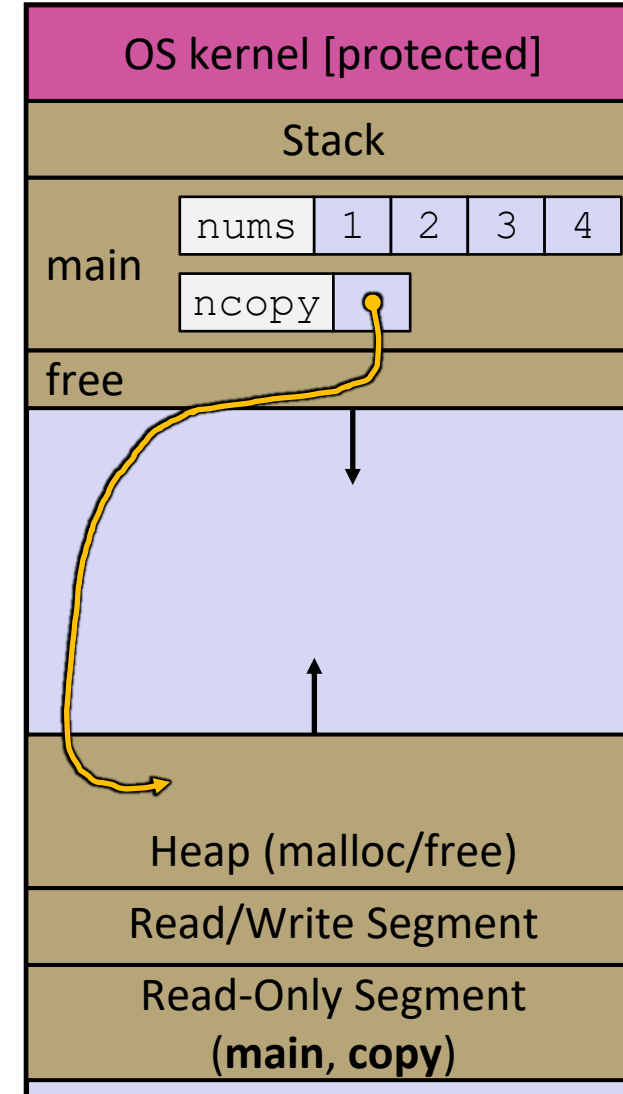
int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Heap and Stack Example (11/11)



arraycopy.c

```
#include <stdlib.h>


int* copy(int a[], int size) {
    int i, *a2;

    a2 = malloc(size*sizeof(int));
    if (a2 == NULL)
        return NULL;

    for (i = 0; i < size; i++)
        a2[i] = a[i];

    return a2;
}

int main(int argc, char** argv) {
    int nums[4] = {1, 2, 3, 4};
    int* ncopy = copy(nums, 4);
    // .. do stuff with the array ..
    free(ncopy);
    return 0;
}
```



# Memory Corruption



- There are all sorts of ways to corrupt memory in C

memcorrupt.c

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char** argv) {
    int a[2];
    int* b = malloc(2*sizeof(int));
    int* c;

    a[2] = 5;    // assign past the end of an array
    b[0] += 2;    // assume malloc zeros out memory
    c = b+3;      // mess up your pointer arithmetic
    free(&(a[0])); // free something not malloc'ed
    free(b);
    free(b);      // double-free the same block
    b[0] = 5;     // use a freed pointer

    // any many more!
    return 0;
}
```

# Memory Leak (1/2)



- A **memory leak** occurs when
  - code fails to deallocate dynamically-allocated memory that is no longer used
  - *e.g.* forget to **free** malloc-ed block, lose/change pointer to malloc-ed block



# Memory Leak (2/2)



- Implication: program's VM footprint will keep growing
  - This might be OK for *short-lived* program, since memory deallocated when program ends
  - Usually has bad repercussions for *long-lived* programs
    - Might slow down over time (*e.g.* lead to VM thrashing)
    - Might exhaust all available memory and crash
    - Other programs might get starved of memory

# Ex: malloc.c



```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    int *p;      // pointer to the dynamically allocated blocks
    int n, i;

    printf("The number of your inputs:> ");
    scanf("%d", &n);
    if (n<=0){ // checking the number
        printf("Error: Wrong numbers.\n");
        printf("Program ended...\n");
        return -1;
    }
    p = (int *) malloc (n*sizeof(int));
    if (p== NULL){
        printf("Error: Not enough memory.\n");
        printf("Program ended...\n");
        return -1;
    }
    for (i=0; i< n; ++i)
        scanf("%d", &p[i]);
    printf("Printing the numbers in reverse order.\n");
    for (i=n-1; i>=0 ; --i)
        printf("%d\t", p[i]);
    printf("\n");
}
```

Allocating memory dynamically to  
store the integer to be entered



# C library macro - assert()



- **void assert(int expression)**

- allows diagnostic information to be written to the standard error file.
- **expression** – This can be a variable or any C expression
  - evaluates to TRUE, assert() does nothing.
  - evaluates to FALSE, assert() displays an error message on **stderr** (standard error stream to display error messages and diagnostics) and aborts program execution.

# Ex: astest.c




```
#include <stdio.h>
#include <string.h>
#include <assert.h>    // assert() is defined

void copy(char *dest, char *src)
{
    assert(dest != NULL);    // if dest==NULL, then abort
    assert(src != NULL);    // if src==NULL, then abort
    strcpy(dest, src);    // copy string
}

int main()
{
    char s1[100];
    char *s2 = "Hello, world!";

    copy(s1, s2);    // normal execution

    copy(NULL, s2);    // src is NULL
    // Assertion failed: dest != NULL,
    return 0;
}
```



```
$ gcc -o astest astest.c
$ ./astest
astest: astest.c:7: copy:
    Assertion `dest != NULL' failed.
Aborted (core dumped)
```

# Q&A

